

Practical Machine Learning Course Project

Melanie Carlson

Sunday, June 21, 2015

The goal of this project is to determine how the participant did the exercise. This report will answer the following: How the model was built, How it was cross validated, Expected Sample Error and why I made the choices I did.

Question

Use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants who were asked to perform barbell lifts correctly and incorrectly in 5 different ways to create an algorithm that correctly identifies the activity quality. The five ways, as described in the study, were “exactly according to the specification (Class A), throwing the elbows to the front (Class B), lifting the dumbbell only halfway (Class C), lowering the dumbbell only halfway (Class D) and throwing the hips to the front (Class E). Class A corresponds to the specified execution of the exercise, while the other 4 classes correspond to common mistakes

Input Data

Load necessary libraries

```
library(AppliedPredictiveModeling)

## Warning: package 'AppliedPredictiveModeling' was built under R version
## 3.2.1

library(caret)

## Warning: package 'caret' was built under R version 3.2.1
## Loading required package: lattice
## Loading required package: ggplot2

library(rattle)

## Warning: package 'rattle' was built under R version 3.2.1
## Rattle: A free graphical interface for data mining with R.
## Version 3.4.1 Copyright (c) 2006-2014 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.

library(rpart.plot)

## Loading required package: rpart

library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.2.1
## randomForest 4.6-10
## Type rfNews() to see new features/changes/bug fixes.
library(knitr)
## Warning: package 'knitr' was built under R version 3.2.1
library(e1071)
## Warning: package 'e1071' was built under R version 3.2.1
```

Upload Data

```
setwd("C:/Users/Melanie/Desktop/R Code Class/Machine Learning")
df_training <- read.csv("pml-training.csv", na.strings=c("NA", ""),
header=TRUE)
colnames_train <- colnames(df_training)
df_testing <- read.csv("pml-testing.csv", na.strings=c("NA", ""), header=TRUE)
colnames_test <- colnames(df_testing)
```

Verify that the column names (excluding classe and problem_id) are identical in the training and test set.

```
all.equal(colnames_train[1:length(colnames_train)-1],
colnames_test[1:length(colnames_train)-1])
## [1] TRUE
```

Features

Remove Incomplete Data

```
# Count the number of non-NAs in each col.
nonNAs <- function(x) {
  as.vector(apply(x, 2, function(x) length(which(!is.na(x)))))
}

# Build vector of missing data or NA columns to drop.
colcnts <- nonNAs(df_training)
```

```

drops <- c()
for (cnt in 1:length(colcnts)) {
  if (colcnts[cnt] < nrow(df_training)) {
    drops <- c(drops, colnames_train[cnt])
  }
}

# Drop NA data and the first 7 columns as they're unnecessary for predicting.
df_training <- df_training[,!(names(df_training) %in% drops)]
df_training <- df_training[,8:length(colnames(df_training))]

df_testing <- df_testing[,!(names(df_testing) %in% drops)]
df_testing <- df_testing[,8:length(colnames(df_testing))]

```

Show remaining columns.

```

colnames(df_training)
##  [1] "roll_belt"          "pitch_belt"         "yaw_belt"
##  [4] "total_accel_belt"   "gyros_belt_x"       "gyros_belt_y"
##  [7] "gyros_belt_z"       "accel_belt_x"       "accel_belt_y"
## [10] "accel_belt_z"       "magnet_belt_x"      "magnet_belt_y"
## [13] "magnet_belt_z"     "roll_arm"           "pitch_arm"
## [16] "yaw_arm"           "total_accel_arm"    "gyros_arm_x"
## [19] "gyros_arm_y"       "gyros_arm_z"       "accel_arm_x"
## [22] "accel_arm_y"       "accel_arm_z"       "magnet_arm_x"
## [25] "magnet_arm_y"      "magnet_arm_z"      "roll_dumbbell"
## [28] "pitch_dumbbell"    "yaw_dumbbell"       "total_accel_dumbbell"
## [31] "gyros_dumbbell_x"  "gyros_dumbbell_y"   "gyros_dumbbell_z"
## [34] "accel_dumbbell_x"  "accel_dumbbell_y"   "accel_dumbbell_z"
## [37] "magnet_dumbbell_x" "magnet_dumbbell_y"  "magnet_dumbbell_z"
## [40] "roll_forearm"      "pitch_forearm"      "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x"    "gyros_forearm_y"
## [46] "gyros_forearm_z"   "accel_forearm_x"    "accel_forearm_y"
## [49] "accel_forearm_z"   "magnet_forearm_x"   "magnet_forearm_y"

```

```
## [52] "magnet_forearm_z"      "classe"
colnames(df_testing)
## [1] "roll_belt"      "pitch_belt"      "yaw_belt"
## [4] "total_accel_belt"  "gyros_belt_x"    "gyros_belt_y"
## [7] "gyros_belt_z"     "accel_belt_x"    "accel_belt_y"
## [10] "accel_belt_z"     "magnet_belt_x"   "magnet_belt_y"
## [13] "magnet_belt_z"    "roll_arm"        "pitch_arm"
## [16] "yaw_arm"          "total_accel_arm" "gyros_arm_x"
## [19] "gyros_arm_y"      "gyros_arm_z"     "accel_arm_x"
## [22] "accel_arm_y"      "accel_arm_z"     "magnet_arm_x"
## [25] "magnet_arm_y"     "magnet_arm_z"    "roll_dumbbell"
## [28] "pitch_dumbbell"   "yaw_dumbbell"    "total_accel_dumbbell"
## [31] "gyros_dumbbell_x" "gyros_dumbbell_y" "gyros_dumbbell_z"
## [34] "accel_dumbbell_x" "accel_dumbbell_y" "accel_dumbbell_z"
## [37] "magnet_dumbbell_x" "magnet_dumbbell_y" "magnet_dumbbell_z"
## [40] "roll_forearm"     "pitch_forearm"   "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x" "gyros_forearm_y"
## [46] "gyros_forearm_z"  "accel_forearm_x" "accel_forearm_y"
## [49] "accel_forearm_z"  "magnet_forearm_x" "magnet_forearm_y"
## [52] "magnet_forearm_z"  "problem_id"
```

Check Covariates for Variability

```
nsv <- nearZeroVar(df_training, saveMetrics=TRUE)
nsv
```

##	freqRatio	percentUnique	zeroVar	nzv
## roll_belt	1.101904	6.7781062	FALSE	FALSE
## pitch_belt	1.036082	9.3772296	FALSE	FALSE
## yaw_belt	1.058480	9.9734991	FALSE	FALSE
## total_accel_belt	1.063160	0.1477933	FALSE	FALSE
## gyros_belt_x	1.058651	0.7134849	FALSE	FALSE
## gyros_belt_y	1.144000	0.3516461	FALSE	FALSE
## gyros_belt_z	1.066214	0.8612782	FALSE	FALSE
## accel_belt_x	1.055412	0.8357966	FALSE	FALSE

## accel_belt_y	1.113725	0.7287738	FALSE	FALSE
## accel_belt_z	1.078767	1.5237998	FALSE	FALSE
## magnet_belt_x	1.090141	1.6664968	FALSE	FALSE
## magnet_belt_y	1.099688	1.5187035	FALSE	FALSE
## magnet_belt_z	1.006369	2.3290184	FALSE	FALSE
## roll_arm	52.338462	13.5256345	FALSE	FALSE
## pitch_arm	87.256410	15.7323412	FALSE	FALSE
## yaw_arm	33.029126	14.6570176	FALSE	FALSE
## total_accel_arm	1.024526	0.3363572	FALSE	FALSE
## gyros_arm_x	1.015504	3.2769341	FALSE	FALSE
## gyros_arm_y	1.454369	1.9162165	FALSE	FALSE
## gyros_arm_z	1.110687	1.2638875	FALSE	FALSE
## accel_arm_x	1.017341	3.9598410	FALSE	FALSE
## accel_arm_y	1.140187	2.7367241	FALSE	FALSE
## accel_arm_z	1.128000	4.0362858	FALSE	FALSE
## magnet_arm_x	1.000000	6.8239731	FALSE	FALSE
## magnet_arm_y	1.056818	4.4439914	FALSE	FALSE
## magnet_arm_z	1.036364	6.4468454	FALSE	FALSE
## roll_dumbbell	1.022388	84.2065029	FALSE	FALSE
## pitch_dumbbell	2.277372	81.7449801	FALSE	FALSE
## yaw_dumbbell	1.132231	83.4828254	FALSE	FALSE
## total_accel_dumbbell	1.072634	0.2191418	FALSE	FALSE
## gyros_dumbbell_x	1.003268	1.2282132	FALSE	FALSE
## gyros_dumbbell_y	1.264957	1.4167771	FALSE	FALSE
## gyros_dumbbell_z	1.060100	1.0498420	FALSE	FALSE
## accel_dumbbell_x	1.018018	2.1659362	FALSE	FALSE
## accel_dumbbell_y	1.053061	2.3748853	FALSE	FALSE
## accel_dumbbell_z	1.133333	2.0894914	FALSE	FALSE
## magnet_dumbbell_x	1.098266	5.7486495	FALSE	FALSE
## magnet_dumbbell_y	1.197740	4.3012945	FALSE	FALSE
## magnet_dumbbell_z	1.020833	3.4451126	FALSE	FALSE
## roll_forearm	11.589286	11.0895933	FALSE	FALSE
## pitch_forearm	65.983051	14.8557741	FALSE	FALSE
## yaw_forearm	15.322835	10.1467740	FALSE	FALSE

## total_accel_forearm	1.128928	0.3567424	FALSE	FALSE
## gyros_forearm_x	1.059273	1.5187035	FALSE	FALSE
## gyros_forearm_y	1.036554	3.7763735	FALSE	FALSE
## gyros_forearm_z	1.122917	1.5645704	FALSE	FALSE
## accel_forearm_x	1.126437	4.0464784	FALSE	FALSE
## accel_forearm_y	1.059406	5.1116094	FALSE	FALSE
## accel_forearm_z	1.006250	2.9558659	FALSE	FALSE
## magnet_forearm_x	1.012346	7.7667924	FALSE	FALSE
## magnet_forearm_y	1.246914	9.5403119	FALSE	FALSE
## magnet_forearm_z	1.000000	8.5771073	FALSE	FALSE
## classe	1.469581	0.0254816	FALSE	FALSE

All False, so no need to remove any covariates due to lack of variability

Algorithm

Devided the data up into 4 sets, so able to do multiple trials of the algorithm to avoid overfitting the predictor and so things would run quicker.

```
set.seed(2)
ids_small <- createDataPartition(y=df_training$classe, p=0.25, list=FALSE)
df_small1 <- df_training[ids_small,]
df_remainder <- df_training[-ids_small,]
set.seed(666)
ids_small <- createDataPartition(y=df_remainder$classe, p=0.33, list=FALSE)
df_small2 <- df_remainder[ids_small,]
df_remainder <- df_remainder[-ids_small,]
set.seed(666)
ids_small <- createDataPartition(y=df_remainder$classe, p=0.5, list=FALSE)
df_small3 <- df_remainder[ids_small,]
df_small4 <- df_remainder[-ids_small,]
# Divide each of these 4 sets into training (60%) and test (40%) sets.
set.seed(666)
inTrain <- createDataPartition(y=df_small1$classe, p=0.6, list=FALSE)
```

```

df_small_training1 <- df_small1[inTrain,]
df_small_testing1 <- df_small1[-inTrain,]
set.seed(666)
inTrain <- createDataPartition(y=df_small2$classe, p=0.6, list=FALSE)
df_small_training2 <- df_small2[inTrain,]
df_small_testing2 <- df_small2[-inTrain,]
set.seed(666)
inTrain <- createDataPartition(y=df_small3$classe, p=0.6, list=FALSE)
df_small_training3 <- df_small3[inTrain,]
df_small_testing3 <- df_small3[-inTrain,]
set.seed(666)
inTrain <- createDataPartition(y=df_small4$classe, p=0.6, list=FALSE)
df_small_training4 <- df_small4[inTrain,]
df_small_testing4 <- df_small4[-inTrain,]

```

Parameters

Used classification trees “out of the box” and then introduce preprocessing and cross validation.

Evaluation

Classification Tree

```

set.seed(2)

##Train

modFit <- train(df_small_training1$classe ~ ., data = df_small_training1,
method="rpart")
print(modFit, digits=3)

## CART
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'

```

```
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
##
## Summary of sample sizes: 2946, 2946, 2946, 2946, 2946, 2946, ...
##
## Resampling results across tuning parameters:
##
##   cp      Accuracy  Kappa  Accuracy SD  Kappa SD
##   0.0356  0.517     0.3773  0.0685      0.1054
##   0.0596  0.375     0.1454  0.0399      0.0645
##   0.1157  0.340     0.0907  0.0373      0.0523
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was cp = 0.0356.
print(modFit$finalModel, digits=3)
## n= 2946
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##  1) root 2946 2110 A (0.28 0.19 0.17 0.16 0.18)
##    2) roll_belt< 130 2696 1860 A (0.31 0.21 0.19 0.18 0.11)
##      4) pitch_forearm< -34 227      1 A (1 0.0044 0 0 0) *
##      5) pitch_forearm>=-34 2469 1860 A (0.25 0.23 0.21 0.2 0.12)
##        10) magnet_dumbbell_y< 426 2051 1460 A (0.29 0.18 0.24 0.19 0.099)
##          20) roll_forearm< 122 1287  752 A (0.42 0.17 0.19 0.16 0.057) *
##          21) roll_forearm>=122 764  513 C (0.075 0.19 0.33 0.24 0.17) *
##        11) magnet_dumbbell_y>=426 418  219 B (0.038 0.48 0.038 0.23 0.22) *
##    3) roll_belt>=130 250      3 E (0.012 0 0 0 0.99) *
fancyRpartPlot(modFit$finalModel)
```

```
##Run Against my Test Set
```



```

predictions <- predict(modFit, newdata=df_small_testing1)
print(confusionMatrix(predictions, df_small_testing1$classe), digits=4)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction   A    B    C    D    E
##      A  515  154  146  133   44
##      B    7  143   15   59   54
##      C   32   83  181  129   97
##      D    0    0    0    0    0
##      E    4    0    0    0  165
##
## Overall Statistics
##
##              Accuracy : 0.512
##              95% CI : (0.4896, 0.5343)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.3632
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          0.9229  0.37632  0.5292  0.0000  0.45833
## Specificity          0.6600  0.91461  0.7894  1.0000  0.99750
## Pos Pred Value       0.5192  0.51439  0.3467      NaN  0.97633
## Neg Pred Value       0.9556  0.85918  0.8881  0.8363  0.89118
## Prevalence           0.2845  0.19378  0.1744  0.1637  0.18358
## Detection Rate       0.2626  0.07292  0.0923  0.0000  0.08414
## Detection Prevalence 0.5059  0.14176  0.2662  0.0000  0.08618
## Balanced Accuracy    0.7915  0.64546  0.6593  0.5000  0.72792

```

Low Accuracy, so attempted again using reprocessing and cross validation

```
set.seed(2)

##Train

modFit <- train(df_small_training1$classe ~ ., preProcess=c("center",
"scale"), trControl=trainControl(method = "cv", number = 4), data =
df_small_training1, method="rpart")

print(modFit, digits=3)

## CART
##
## 2946 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (4 fold)
##
## Summary of sample sizes: 2209, 2211, 2209, 2209
##
## Resampling results across tuning parameters:
##
##    cp          Accuracy  Kappa  Accuracy SD  Kappa SD
##    0.0356   0.525      0.3954  0.01009      0.0138
##    0.0596   0.367      0.1265  0.00378      0.0054
##    0.1157   0.345      0.0932  0.04061      0.0622
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was cp = 0.0356.

##Run Against my Test Set

predictions <- predict(modFit, newdata=df_small_testing1)
print(confusionMatrix(predictions, df_small_testing1$classe), digits=4)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    A    B    C    D    E
```

```

##          A 515 154 146 133  44
##          B   7 143  15  59  54
##          C  32  83 181 129  97
##          D   0   0   0   0   0
##          E   4   0   0   0 165
##
## Overall Statistics
##
##          Accuracy : 0.512
##          95% CI : (0.4896, 0.5343)
##    No Information Rate : 0.2845
##    P-Value [Acc > NIR] : < 2.2e-16
##
##          Kappa : 0.3632
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##          Class: A Class: B Class: C Class: D Class: E
## Sensitivity          0.9229  0.37632  0.5292  0.0000  0.45833
## Specificity          0.6600  0.91461  0.7894  1.0000  0.99750
## Pos Pred Value       0.5192  0.51439  0.3467      NaN  0.97633
## Neg Pred Value       0.9556  0.85918  0.8881  0.8363  0.89118
## Prevalence           0.2845  0.19378  0.1744  0.1637  0.18358
## Detection Rate       0.2626  0.07292  0.0923  0.0000  0.08414
## Detection Prevalence 0.5059  0.14176  0.2662  0.0000  0.08618
## Balanced Accuracy     0.7915  0.64546  0.6593  0.5000  0.72792

```

Little improvement on accuracy, so decided to use Random Forest.

Random Forest

```

set.seed(2)

##Train

```

```
modFit <- train(df_small_training1$classe ~ ., method="rf",
trControl=trainControl(method = "cv", number = 4), data=df_small_training1)
print(modFit, digits=3)
```

```
## Random Forest
##
## 2946 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
##
## Summary of sample sizes: 2209, 2211, 2209, 2209
##
## Resampling results across tuning parameters:
##
##    mtry  Accuracy  Kappa  Accuracy SD  Kappa SD
##    2     0.946     0.931  0.01101      0.01395
##   27     0.952     0.940  0.00573      0.00725
##   52     0.946     0.932  0.00806      0.01022
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was mtry = 27.
```

```
##Run Against my Test Set
```

```
predictions <- predict(modFit, newdata=df_small_testing1)
print(confusionMatrix(predictions, df_small_testing1$classe), digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
##              Reference
## Prediction   A    B    C    D    E
##      A 553    8    0    2    0
##      B   2 361   18    2    9
##      C   2  10 320   14    1
##      D   0   1   4 301    1
```

```
##           E    1    0    0    2 349
##
## Overall Statistics
##
##           Accuracy : 0.9607
##           95% CI : (0.9512, 0.9689)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9503
##           McNemar's Test P-Value : 0.002504
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.9910   0.9500   0.9357   0.9377   0.9694
## Specificity           0.9929   0.9804   0.9833   0.9963   0.9981
## Pos Pred Value        0.9822   0.9209   0.9222   0.9805   0.9915
## Neg Pred Value        0.9964   0.9879   0.9864   0.9879   0.9932
## Prevalence            0.2845   0.1938   0.1744   0.1637   0.1836
## Detection Rate        0.2820   0.1841   0.1632   0.1535   0.1780
## Detection Prevalence  0.2871   0.1999   0.1770   0.1566   0.1795
## Balanced Accuracy      0.9920   0.9652   0.9595   0.9670   0.9838
##Ran Against Course Provided Test Set
print(predict(modFit, newdata=df_testing))
## [1] B A B A A E D D A A B C B A E E A B B B
## Levels: A B C D E
```

attempted again using reprocessing and cross validation

```
set.seed(2)

##Train

modFit <- train(df_small_training1$classe ~ ., method="rf",
preProcess=c("center", "scale"), trControl=trainControl(method = "cv", number
= 4), data=df_small_training1)

print(modFit, digits=3)
```

```
## Random Forest
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (4 fold)
##
## Summary of sample sizes: 2209, 2211, 2209, 2209
##
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.946 0.931 0.01276 0.01616
## 27 0.951 0.938 0.00588 0.00744
## 52 0.946 0.932 0.00798 0.01014
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

```
##Run Against my Test Set
```

```
predictions <- predict(modFit, newdata=df_small_testing1)
print(confusionMatrix(predictions, df_small_testing1$classe), digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  A  B  C  D  E
```

```
##           A 553  7  0  2  0
```

```
##           B  3 362 17  2  7
```

```
##           C  2 10 321 12  1
```

```
##           D  0  1  3 303  1
```

```
##           E  0  0  1  2 351
```

```
##
```

```
## Overall Statistics
```

```
##
##              Accuracy : 0.9638
##              95% CI : (0.9545, 0.9716)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9542
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          0.9910   0.9526   0.9386   0.9439   0.9750
## Specificity          0.9936   0.9817   0.9846   0.9970   0.9981
## Pos Pred Value       0.9840   0.9258   0.9277   0.9838   0.9915
## Neg Pred Value       0.9964   0.9885   0.9870   0.9891   0.9944
## Prevalence           0.2845   0.1938   0.1744   0.1637   0.1836
## Detection Rate       0.2820   0.1846   0.1637   0.1545   0.1790
## Detection Prevalence 0.2866   0.1994   0.1764   0.1571   0.1805
## Balanced Accuracy     0.9923   0.9671   0.9616   0.9704   0.9866
##
## Ran Against Course Provided Test Set
print(predict(modFit, newdata=df_testing))
##      [1] B A B A A E D B A A B C B A E E A B B B
##
## Levels: A B C D E
```

Mixed results to decided to run again using my second set of data with only Cross Validation

```
set.seed(2)

## Train
modFit <- train(df_small_training2$classe ~ ., method="rf",
preProcess=c("center", "scale"), trControl=trainControl(method = "cv", number
= 4), data=df_small_training2)

print(modFit, digits=3)

## Random Forest
##
```

```
## 2917 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (4 fold)
##
## Summary of sample sizes: 2188, 2188, 2188, 2187
##
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.948 0.934 0.00783 0.00991
## 27 0.955 0.943 0.00847 0.01069
## 52 0.948 0.934 0.00605 0.00760
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

```
##Run Against my Test Set
```

```
predictions <- predict(modFit, newdata=df_small_testing2)
print(confusionMatrix(predictions, df_small_testing2$classe), digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  A  B  C  D  E
```

```
##           A 549 10  0  0  0
```

```
##           B  2 363  8  1  0
```

```
##           C  0  3 327 18  3
```

```
##           D  0  0  3 298  6
```

```
##           E  1  0  0  1 348
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.9711
```



```
##              95% CI : (0.9627, 0.9781)
##      No Information Rate : 0.2844
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9635
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          0.9946   0.9654   0.9675   0.9371   0.9748
## Specificity          0.9928   0.9930   0.9850   0.9945   0.9987
## Pos Pred Value       0.9821   0.9706   0.9316   0.9707   0.9943
## Neg Pred Value       0.9978   0.9917   0.9931   0.9878   0.9943
## Prevalence           0.2844   0.1937   0.1741   0.1638   0.1839
## Detection Rate       0.2828   0.1870   0.1685   0.1535   0.1793
## Detection Prevalence 0.2880   0.1927   0.1808   0.1582   0.1803
## Balanced Accuracy     0.9937   0.9792   0.9762   0.9658   0.9868
##Ran Against Course Provided Test Set
print(predict(modFit, newdata=df_testing))
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Looked good, so ran my 3rd set of data with only Cross Validation

```
set.seed(2)
##Train
modFit <- train(df_small_training3$classe ~ ., method="rf",
preProcess=c("center", "scale"), trControl=trainControl(method = "cv", number
= 4), data=df_small_training3)
print(modFit, digits=3)
## Random Forest
##
## 2960 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
```

```
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (4 fold)
##
## Summary of sample sizes: 2220, 2220, 2220, 2220
##
## Resampling results across tuning parameters:
##
##   mtry  Accuracy  Kappa  Accuracy SD  Kappa SD
##    2    0.951    0.938  0.00943    0.0120
##   27    0.948    0.934  0.01547    0.0197
##   52    0.942    0.927  0.01777    0.0225
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was mtry = 2.
```

```
##Run Against my Test Set
```

```
predictions <- predict(modFit, newdata=df_small_testing3)
print(confusionMatrix(predictions, df_small_testing3$classe), digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
## Prediction  A   B   C   D   E
##           A 558  16   0   1   1
##           B   1 357  23   0   1
##           C   0   7 319  24   5
##           D   1   1   2 298   3
##           E   0   0   0   0 352
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.9563
##           95% CI   : (0.9464, 0.9649)
##           No Information Rate : 0.2843
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
##          Kappa : 0.9447
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##          Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9964   0.9370   0.9273   0.9226   0.9724
## Specificity      0.9872   0.9843   0.9779   0.9957   1.0000
## Pos Pred Value   0.9688   0.9346   0.8986   0.9770   1.0000
## Neg Pred Value   0.9986   0.9849   0.9845   0.9850   0.9938
## Prevalence       0.2843   0.1934   0.1746   0.1640   0.1838
## Detection Rate   0.2832   0.1812   0.1619   0.1513   0.1787
## Detection Prevalence 0.2924   0.1939   0.1802   0.1548   0.1787
## Balanced Accuracy 0.9918   0.9606   0.9526   0.9592   0.9862
##Ran Against Course Provided Test Set
print(predict(modFit, newdata=df_testing))
##  [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Looked good, so ran my final set of data with only Cross Validation

```
set.seed(2)
##Train
modFit <- train(df_small_training4$classe ~ ., method="rf",
preProcess=c("center", "scale"), trControl=trainControl(method = "cv", number
= 4), data=df_small_training4)
print(modFit, digits=3)
## Random Forest
##
## 2958 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (4 fold)
```

```
##
## Summary of sample sizes: 2217, 2219, 2219, 2219
##
## Resampling results across tuning parameters:
##
##   mtry  Accuracy  Kappa  Accuracy SD  Kappa SD
##    2    0.949    0.936  0.00633    0.00805
##   27    0.955    0.943  0.00697    0.00886
##   52    0.950    0.937  0.00403    0.00514
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was mtry = 27.
```

```
##Run Against my Test Set
```

```
predictions <- predict(modFit, newdata=df_small_testing4)
print(confusionMatrix(predictions, df_small_testing4$classe), digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction  A   B   C   D   E
##           A 556  20   2   0   0
##           B   3 341   8   0   3
##           C   0  19 329   4   3
##           D   1   1   4 318   4
##           E   0   0   0   1 352
##
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.9629
##           95% CI : (0.9536, 0.9708)
##           No Information Rate : 0.2844
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9531
##           McNemar's Test P-Value : NA
```

```
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9929   0.8950   0.9592   0.9845   0.9724
## Specificity      0.9844   0.9912   0.9840   0.9939   0.9994
## Pos Pred Value   0.9619   0.9606   0.9268   0.9695   0.9972
## Neg Pred Value   0.9971   0.9752   0.9913   0.9970   0.9938
## Prevalence       0.2844   0.1935   0.1742   0.1640   0.1838
## Detection Rate   0.2824   0.1732   0.1671   0.1615   0.1788
## Detection Prevalence 0.2936  0.1803  0.1803  0.1666  0.1793
## Balanced Accuracy 0.9886   0.9431   0.9716   0.9892   0.9859
##Ran Against Course Provided Test Set
print(predict(modFit, newdata=df_testing))
##  [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Out of Sample Error Rate

Average of the sample error rates derived by applying the random forest method with both preprocessing and cross validation against test sets 1-4 yielding a predicted out of sample rate of 0.03584.